

42. Unique ID Generator

Requirements clarification

Functional requirements

- Mandatory requirements
 - ID must be unique among the whole distributed system.
- Optional requirements
 - IDs can be ordered by time.

Non-functional requirements


- The generator should be highly available.
- ID generation should happen in real-time with minimal latency.

High-level design

Architecture options


Centralized: Use one centralized service to generate ID.

- ◆ Pros
 - Easy to implement.
- ◆ Cons
 - Introduce single point of failure.



Decentralized: Each server generates ID respectively.


- Pros
 - Improve availability.
- Cons



Detailed design

Solution options for centralized architecture

- Zookeeper
 - Concept
 - ◆ Distributed coordinator to give each server a unique unused range of IDs.



- Pros
 - ◆ Easy to implement.
- Cons
 - ◆ IDs cannot be ordered by time cross multiple servers.
 - ◆ Need to consider the size of each range.


Solution options for decentralized architecture

UUID

- Concepts
 - UUID consists of 32 hexadecimal (base-16) digits, displayed in five groups separated by hyphens.
 - 123e4567-e89b-12d3-a456-426614174000
 - 128-bit.
- Pros
 - Very low probability of getting collision.
 - The ID generators in the services don't need to coordinate each other.
- Cons
 - IDs cannot be ordered by time.

Snowflake ID

- Concepts
 - The format was created by Twitter and is used for the IDs of tweets.
 - Divide an ID into different sections:
 - ◆ Sign (1 bit): Always be 0 (Reserved for future used).
 - ◆ Timestamp (41 bits): Milliseconds since the epoch or custom epoch.
 - ◆ Machine ID (10 bits): Hold up to 1024 machines.
 - ◆ Sequence number (12 bits): The sequence number is incremented by 1 and is reset to 0 every millisecond.
 - ◆ 64-bit.
 - ◆ Each machine can generate IDs by itself.



- Pros
 - IDs can be ordered by time.
 - The binary representation of the timestamp field can be converted to/from a real date and time.