

44. URL Shortening System

Requirements clarification

Functional requirements

- URL Shortening (Write): Given an original URL, our service should generate a shorter and unique URL of it.
- URL Redirection (Read): When users access a short URL, our service should redirect them to the original URL.

Optional function requirements

- URL Customization: Users should optionally be able to pick a custom short URL for their original URL.
- URL Expiration: Shorter URL will expire after a standard default timespan. Users should be able to specify the expiration time.

Non-functional requirements

- The system should be highly available (If our service is down, all the URL redirections will start failing).
- URL redirection should happen in real-time with minimal latency.
- Shortened links should not be guessable (not predictable).

Estimation

Traffic estimation

Our system will be read-heavy (Lots of redirection requests compared to new URL shortenings).

- Read-write ratio is 100 : 1 (Assumed)
- Number of read actions and write actions per month

- Number of writes (URL Shortening) per month = 500 millions (Assumed)
- Number of reads (URL Redirection) per month = 500 millions x 100 = 50 billion
- Frequency of read actions and write actions per second (QPS)
- Frequency of writes per second = $500 \text{ millions} / (30 \text{ days} \times 24 \text{ hours} \times 3600 \text{ seconds}) = 200 \text{ times/s}$
- Frequency of reads per second = $200 \text{ times/s} \times 100 = 20000 \text{ times/s}$

Storage estimation

Types

Data: Yes

File: No

Capacity

- Time length of storing a record = 5 years (Assumed)
- Number of records created in 5 years = Number of writes per month x Number of months = $500 \text{ million} \times 5 \text{ years} \times 12 \text{ months} = 30 \text{ billion}$
- Size of one record = 500 bytes (Assumed)
- Total capacity needed in 5 years = $30 \text{ billion} \times 500 \text{ bytes} = 15 \text{ TB}$

Bandwidth estimation

- Write bandwidth = Frequency of writes per second x Size of one record = $200 \text{ times/s} \times 500 \text{ bytes} = 100 \text{ KB/s}$
- Read bandwidth = Frequency of reads per second x Size of one record = $20000 \text{ times/s} \times 500 \text{ bytes} = 10 \text{ MB/s}$

System interface definition

Interface 1

- **createUrl(original_url)**
 - Function
 - ◆ Create a new shorter URL.
 - Parameters
 - ◆ original_url (string): Original URL to be shortened.
 - Return
 - ◆ The short URL.

Interface 2

- **getURL(api_key, short_url)**
 - Function
 - ◆ Get the original long URL of a short URL.
 - Parameters
 - ◆ short_url (string): The short URL to be redirected.
 - Return
 - ◆ The original long URL.

Data model definition

Schema

- Table 1: URL

- Description

- ◆ Store URL mappings.

- Columns

Column Name	Column Type	PK	Description
ID	int	PK	
ShortUrl	string		The short URL.
LongUrl	string		The original long URL.

Database

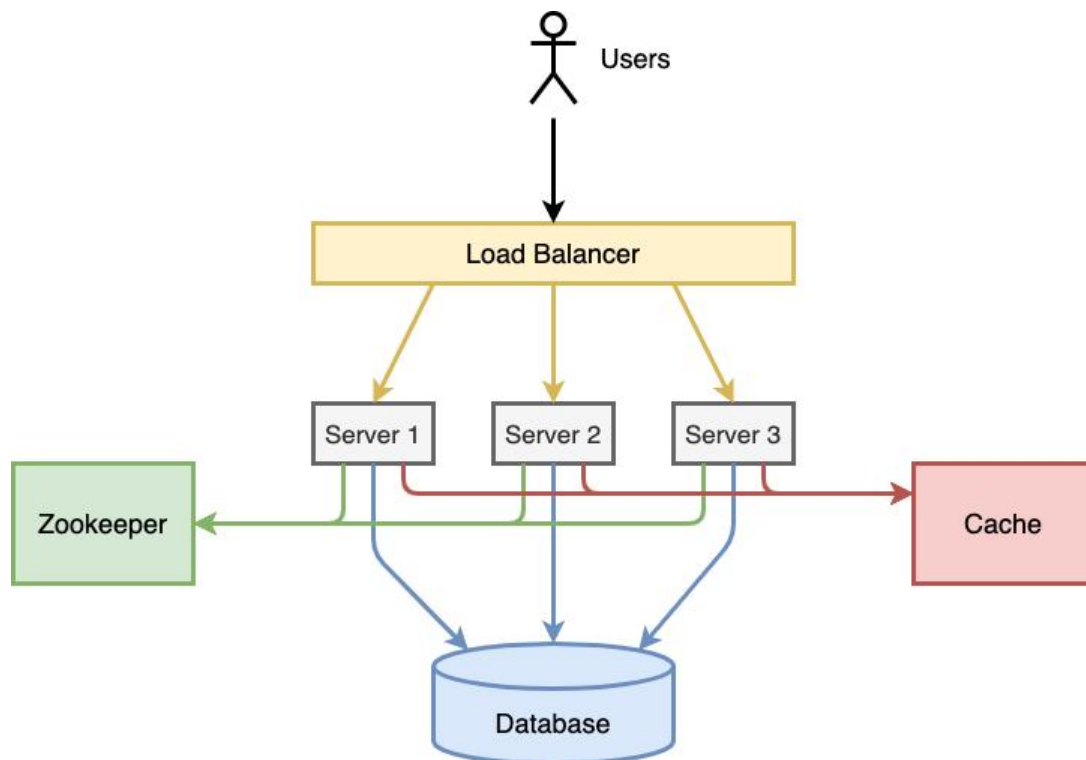
- NoSQL

- Reason

- ◆ No relation need to look up.

- ◆ NoSQL is good at scaling.

High-level design



Zookeeper

- Distributed coordinator to give each server a unique unused range of keys.

Cache

- Stores the top 20% most used URLs.
- When a server receives a URL query request, it can search the cache first. If the target URL is in the cache, it can query the database.


Database

- Stores URLs and users.

Detailed design

URL shortening

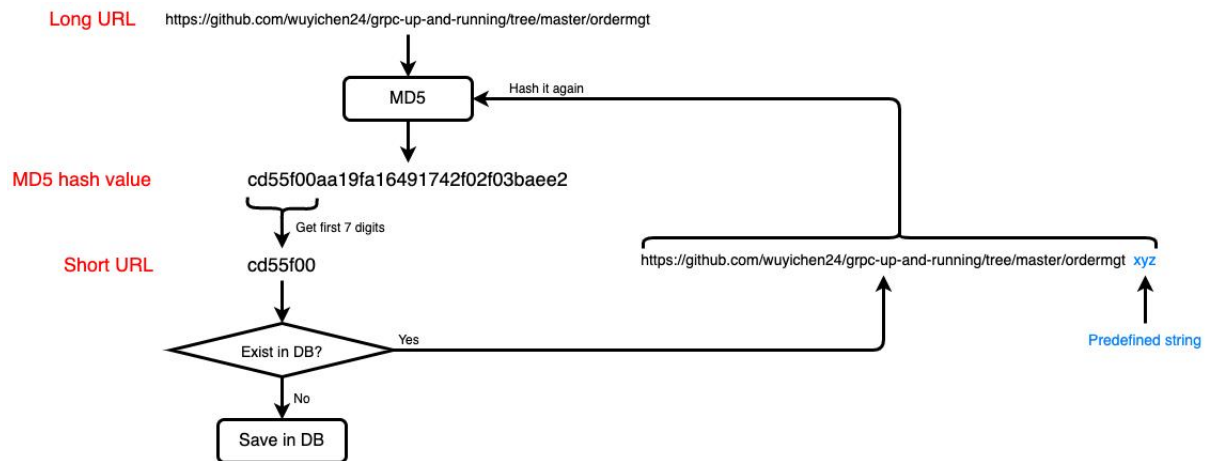
Process



Choices of hash function

Use existing hashing algorithm with collision resolution

- ◆ A hash value from an existing hashing algorithm (CRC32, MD5, SHA-1, SHA-2, etc.) is too long, so we cannot use it directly. Our solution is to only use the first 7 characters of a hash value from an existing hashing algorithm.
- ◆ Using the first 7 characters can result in a hash collision more easily.
- ◆ If the first 7 characters has a hash collision, recursively append a new predefined string to the long URL and hash the new long URL again, until no hash collision.



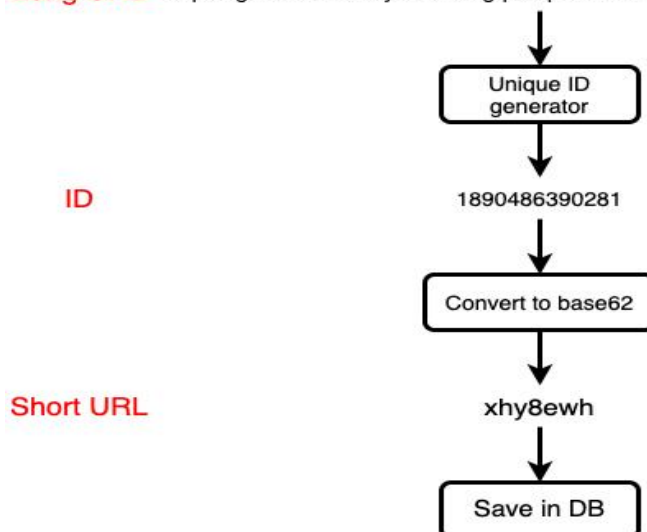
Use base62 conversion

- Convert the unique ID (numeric value) of the new row for the URL mapping table from base 10 to base 62.

Example

- The long URL is `https://github.com/wuyichen24/grpc-up-and-running/tree/master/ordermgt`.
- The new ID for the new row is 1890486390281.
- Convert the ID from base 10 to base 62: xhy8ewh.
- The short URL will be <https://abc.com/xhy8ewh>.

Long URL `https://github.com/wuyichen24/grpc-up-and-running/tree/master/ordermgt`



Uniqueness of short URLs

- **Factors**
 - Number of all possible characters in one digit.
 - The length of a short URL (Number of digits).
- **Calculation:** Number of unique URLs = Number of all possible characters in one digit \times Number of digits
- **Evaluation tradeoffs**
 - Keep short URL as short as possible.
 - Don't let unique short URLs run out easily (Maximal number of URLs > Total number of short URLs created in 5 years).

Solutions

Number of all possible characters in one digit	Length of URLs	Maximal number of URLs
Only numbers (0-9) = 10	7	$10^7 = 10$ million
Base36 ([0-9, a-z]) = 36	7	$36^7 = 78$ billion
Base62 ([0-9, a-z, A-Z]) = 62	7	$62^7 = 3.5$ trillion

Key points

Use Zookeeper as the distributed coordinator to solve the key conflict problem among multiple servers.